

Introduction to Computer Science

Linear Algebra – Arrays

鄒年棣 (Nien-Ti Tsou)

Introduction

- Linear algebra is one of the essential building blocks of computational mathematics. The objects of linear algebra are **vectors and matrices**. The package `NumPy` includes all the necessary tools to manipulate those objects.
- The first task is to build matrices and vectors, or to alter existing ones by slicing.
- The other main task is the `dot` operation, which embodies most of the linear algebra operations (scalar product, matrix-vector product, and matrix-matrix product).
- Finally, various methods are available to solve linear problems.

Vectors

- Creating vectors is as simple as using the function `array` to convert a list to an array:

```
import numpy as np
v = np.array([1., 2., 3.])
```

- The object `v` is now a vector. Here are some illustrations of the basic linear algebra operations on vectors:

```
v1 = np.array([1., 2., 3.])
v2 = np.array([2, 0, 1])
print(v1/2)
print(3*v1 + 2*v2)
```

```
print(np.dot(v1, v2))
print(v1 @ v2) # alternative formulation
print(np.cross(v1, v2))
```

v1		(3,)	[1. 2. 3.]
v2		(3,)	[2 0 1]

```
[0.5  1.   1.5]
[ 7.   6.  11.]
5.0
5.0
[ 2.   5.  -4.]
```

Universal functions: elementwise

- Note that all basic arithmetic operations are performed **elementwise**. Thus, they have an output array that has **the same shape** as the input array.

v1	float64	(3,)	[1. 2. 3.]
v2	int64	(3,)	[2 0 1]

```
print(v1 * v2)
print(v2 / v1)
print(v1 - v2)
print(v1 + v2)
```

```
[
[
[-1.  2.  2.]
[3.  2.  4.]
]
```

- Some operators act elementwise on arrays as well:

```
print(np.cos(np.array([0, (np.pi)/2, np.pi])))
print(np.array([1, 2])**2)
print(2**np.array([1, 2]))
print(np.array([1, 2])**np.array([1, 2]))
```

```
[ 1.0000000e+00  6.123234e-17 -1.0000000e+00]
[1 4]
[2 4]
[ ]
```

Matrix

- A matrix is created in a similar way to a vector, but from a list of lists instead:

```
M = np.array([[1., 2], [0., 1]])  
print(M)
```

```
[[1. 2.]  
 [0. 1.]
```

- Python allows a line break for array creation, which makes it more pleasing to the human eye:

```
M = np.array([[1., 2],  
              [0., 1]])  
print(M)
```

More readable!

```
[[1. 2.]  
 [0. 1.]
```

- The n vector, an $n \times 1$, and a $1 \times n$ matrix are three **different objects** even if they contain the same data.

```
V = np.array([1., 2., 1.])  
R = np.array([[1., 2., 1.]])  
print(V, np.shape(V), np.ndim(V))  
print(R, np.shape(R), np.ndim(R))
```

```
[1. 2. 1.] (3,) 1
```

Tuple!

Array properties

- Arrays are used to manipulate vectors, matrices, and more general **tensors** in NumPy.

```
A = np.array([[1, 2, 3], [3, 4, 6]])  
print(A.shape)  
print(A.dtype)  
print(A.strides)
```

```
int64  
(24, 8)
```

- Type `'int64'` means they use **64 bits or 8 bytes** in memory. The complete array is stored in memory row-wise.
- The distance from `A[0, 0]` to the first element in the next row `A[1, 0]` is thus 24 bytes (three matrix elements) in memory. Correspondingly, the distance in memory between `A[0, 0]` and `A[0, 1]` is 8 bytes (one matrix element). These values are stored in the attribute `strides`.

8 bytes 8 8
→ → →
[[1, 2, 3], [3, 4, 6]]

8 bytes
→
24 bytes ↓
[[1 2 3]
 [3 4 6]]

Array types

- A proper array creation should be like:

```
Vf = np.array([1, 2, 1], dtype=float)
Vc = np.array([1, 2, 1], dtype=complex)
```

Vc	complex128	(3,)	[1.+0.j 2.+0.j 1.+0.j]
Vf	float64	(3,)	[1. 2. 1.]

- When no type is specified, the type is **guessed**. The `array` function chooses the type that allows storing of all the specified values:

```
V = np.array([1, 2])
print(V.dtype)
V = np.array([1., 2])
print(V.dtype)
V = np.array([1. + 0j, 2.])
print(V.dtype)
```

int64
float64
complex128

- Silent type conversion, which might give unexpected results:

```
a = np.array([1, 2, 3])
a[0] = 0.5
print(a)
```



Functions to construct arrays

- The usual way to set up an array is via a list.
- There are also a couple of convenient methods for generating special arrays.

```
r = np.random.rand(2,3)
print(r)

a = np.arange(9.2,20.2,2.2)
b = np.linspace(9.2,20.2,6)
print(a)
print(b)

print(np.zeros((2,3)))
np.zeros(np.shape(r)) # Same as above
print(np.ones((2,3)))
print(np.identity(3))

x = np.diag([4,5,6]) # Generate
print(x)
print(np.diag(x)) # Get values
print(np.diag(x,1)) # Get values
print(np.diag(x,2)) # Get values
```

Note!!! Tuple!



```
[[0.70173912 0.52768381 0.97475031]
 [0.48976793 0.87271617 0.29921848]]
[ 9.2 11.4 13.6 15.8 18.
  9.2 11.4 13.6 15.8 18.]
```

```
[[0. 0. 0.]
 [0. 0. 0.]
 [1. 1. 1.]
 [1. 1. 1.]]
```

```
[[1. 0. 0.]
 [0. 1. 0.]
 [0. 0. 1.]]
```

```
[[4 0 0]
 [0 5 0]
 [0 0 6]]
[
]
[
]
[
]
```


Shaping of Array

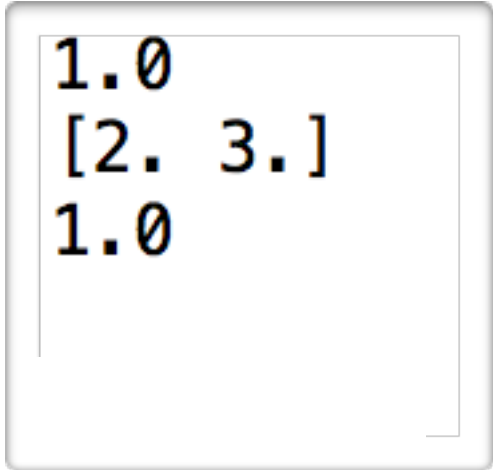
Indexing and slices

- Array entries are accessed by indexes.
 - One pair of brackets:
 - Two pairs of brackets:


```
v = np.array([1., 2., 3])
M = np.array([[1., 2],[3., 4]])
print(v[0]) # works as for lists
print(v[1:])

print(M[0, 0])
print(M[1:]) # returns the matrix array
print(M[1]) # returns the vector array
```

```
v[0] = 10
print(v)
v[:2] = [0, 1]
print(v)
v[:2] = [1, 2, 3]
```



```
1.0
```



```
[10. 2. 3.]
```

Indexing and slices

M[: :]	0	1	2	3
	4	5	6	7
	8	9	10	11

M[: :]	0	1	2	3
	4	5	6	7
	8	9	10	11

M[: :]	0	1	2	3
M[: :]	4	5	6	7
	8	9	10	11

M[: :]	0	1	2	3
	4	5	6	7
	8	9	10	11

- Omitting a dimension: If you omit an index or a slice, NumPy assumes you are taking rows only.

```
M = np.array([[0,1,2,3],[4,5,6,7],[8,9,10,11]])  
print(M[2])
```

```
[ 8  9 10 11]
```

Altering an array using slices

```
M = np.array([[0,1.,2.],[3.,4.,5.],[6,7,8],[9,10,11]])
print(M)
M[1, 2] = 2.0
print(M)
M[2, :] = [1., 2., 3.]
print(M)
M[1:3, :] = np.array([[1., 2., 3.],[-1.,-2., -3.]])
print(M)
```

- There is a distinction between a column matrix and a vector.

```
M[1:4, 2:3] = np.array([[1.], [0.], [-1.0]])
print(M)
M[1:4, 2:3] = np.array([1., 0., -1.0])
```

```
[ [ 0.   1.   2.]  
  [ 3.   4.   5.]  
  [ 6.   7.   8.]  
  [ 9.  10.  11.]]  
  
[ [  
  [  
    [  
      [ ]]  
    ]]  
  [  
    [  
      [ ]]  
    ]]  
  [  
    [  
      [ ]]  
    ]]  
  [  
    [  
      [ ]]  
    ]]
```

$$\begin{bmatrix} \begin{bmatrix} \\ \\ \\ \end{bmatrix} \\ \begin{bmatrix} \\ \\ \\ \end{bmatrix} \\ \begin{bmatrix} \\ \\ \\ \end{bmatrix} \\ \begin{bmatrix} \\ \\ \\ \end{bmatrix} \end{bmatrix}$$

```
ValueError: could not broadcast input array
from shape (3) into shape (3,1)
```

Reshape


- Vector and row matrix can **be viewed as** a column matrix by the method `reshape`:

```
V = np.array([1., 2., 1.])
R = np.array([[1., 2., 1.]])
V2C = V.reshape(3, 1)
R2C = R.reshape(3, 1)
print(V2C, np.shape(V2C))
print(R2C, np.shape(R2C))
```

```
A = np.arange(6)
print(A)
```

```
[[1.]
 [2.]
 [1.]] (3, 1)
[[1.]
 [2.]
 [1.]] (3, 1)
```

```
[0 1 2 3 4 5]
```



0	1	2
3	4	5

`A.reshape(2, 3)`

0	1
2	3
4	5

`A.reshape(3, 2)`

- It is convenient to specify only one shape parameter by setting the free shape parameter `-1`:

```
v = np.array([1, 2, 3, 4, 5, 6, 7, 8])
M = v.reshape(2, -1)
print(M)
M = v.reshape(-1, 2)
print(M)
M = v.reshape(3, -1)
```

```
[[1 2 3 4]
 [5 6 7 8]]
[[1 2]
 [3 4]
 [5 6]
 [7 8]]
```

```
ValueError: cannot reshape array of
size 8 into shape (3,newaxis)
```

Transpose $B_{ij} = A_{ji}$

- A special form of reshaping is transposing. It just switches the two shape elements of the matrix.

In previous page

M

```
[[1, 2],  
 [3, 4],  
 [5, 6],  
 [7, 8]]
```

```
print(M.T)
```

```
[[1 3 5 7]  
 [2 4 6 8]]
```

- Transpose does not copy! Transposition is very similar to reshaping and just returns **a view** on the same array.
- Transposing a vector makes no sense in NumPy:

```
v = np.array([1., 2., 3.])  
print(v)  
print(v.T)
```

```
[1. 2. 3.]  
[1. 2. 3.]
```

- Transpose a vector is probably to create a row or column **matrix**. This is done using `reshape`:

```
print(v.reshape(-1, 1))  
print(v.reshape(1, -1))
```

```
[[1.]  
 [2.]  
 [3.]]  
[[1. 2. 3.]]
```


Array view


- Basic slices (`:`, `T`, `reshape`) return **views**. It is possible to access the object that owns the data using the array attribute `base`. If an array owns its data, the attribute `base` is `None`.
- Changing the elements of a slice of view affects the entire array

```
v = np.array([1., 2., 3.])
print(v)
v1 = v[:2]
print(v1)
v1[0] = 0.
print(v)
print(v1.base is v)
print(v.base is None)
print(id(v), id(v1), v1 is v)
```

```
[1. 2. 3.]
[1. 2.]
[0. 2. 3.]
```

```
17692163088 26283434512
```

- We can say: the data pointed by `v1` is based on (or getting values from) the data pointed by `v`.



Array copy

- Sometimes it is necessary to explicitly request that the data be copied. This is simply achieved with the function `array`:


```
M = np.array([[1.,2.],[3.,4.]])
M2 = np.array(M) # copy of M
M3 = M
N = np.array(M.T) # copy of M.T
N2 = M.T
M[0,0]=10
print(M)
print(M2)
print(M3)
print(N)
print(N2)
```

Alternatively:
`M2 = M.copy()`

- We may verify that the data has indeed been copied:

```
print(M.base is None)
print(M2.base is None)
print(M3.base is M)
print(N.base is None)
print(N2.base is M)
```

```
M4 = M[:]
print(M4.base is M)
```



```
True
True

True
True
```

Why????!!!




Why does `M3.base is M` return `False`?

```
16 print(id(M))
17 print(id(M.base))
18 print(id(M3))
19 print(id(M3.base))
20 print(M3 is M)
21 print(M3.base is M)
22 print(id(N2))
23 print(id(M.T))
24 print(id(M.T.base))
25 print(id(N2.base))
```

1949924536848
1564976272
1949924536848
1564976272
True
False
1949924565760
1949924534272
1949924536848
1949924536848

`M3 is M`

`M3` is not an object extended from `M`, so no base relationship between `M3` and `M`. On the other hand, `N2` gets its values from the memory pointed by `M`.



Thanks to
0611534 陳怡儒

Stacking

- The method to build matrices from a couple of (matching) submatrices is `concatenate`.
- This command stacks the submatrices vertically when `axis=0`. With the `axis=1`, they are stacked horizontally.
- `concatenate` creates a new matrix

```
Z = np.zeros((3,3))
I = np.identity(3)
print(Z)
print(I)
print(np.concatenate((Z,I)))
print(np.concatenate((Z,I),axis = 1))
X = np.concatenate((Z,I),axis = 1)
Z[0,0]=100
print(Z)
print(X)
```

```
[[0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]]
[[1. 0. 0.]
 [0. 1. 0.]
 [0. 0. 1.]]
[[0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]
 [1. 0. 0.]
 [0. 1. 0.]
 [0. 0. 1.]]
[[0. 0. 0. 1. 0. 0.]
 [0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 1.]]
[[100. 0. 0.]
 [ 0. 0. 0.]
 [ 0. 0. 0.]]
[[          1. 0. 0.]
 [          0. 1. 0.]
 [          0. 0. 1.]]
```

More stacking

`hstack`: Used to stack matrices horizontally

`vstack`: Used to stack matrices vertically

`column_stack`: Used to stack vectors in columns

- One may stack vectors row-wise or column-wise using `vstack` and `column_stack`.

```
a=np.array([[1],[2],[3]])  
aaa = np.hstack((a,a,a))  
a[0,0] = 10  
print(a)  
print(aaa)
```

They can be `()` or `[]`

[[1	0]		
	[2]		
	[3]		
	[[1	1	1]
		[2	2	2]
		[3	3	3]

```
v1 = np.array([1,2])  
v2 = np.array([3,4])  
vv = np.vstack([v1,v2])  
print(vv)  
v1v2 = np.column_stack([v1,v2])  
print(v1v2)
```

1	2
3	4

`vstack([v1,v2])`

1	3
2	4

`column_stack([v1,v2])`

Operation on Arrays

Array functions

- Some functions acting on the whole matrix, row-wise, or column-wise, such as `max`, `min`, and `sum`.

```
A = (np.arange(1,9)).reshape(2,-1)
print(A)
print(np.sum(A), np.min(A), np.max(A))
```

```
[[1 2 3 4]
 [5 6 7 8]]
36 1 8
```

- Command `sum` has an optional parameter, `axis`. It allows us to choose along which axis to perform the operation.

```
print(np.sum(A, axis=0))
print(np.sum(A, axis=1))
```

```
[ 6  8 10 12]
[10 26]
```

Comparing arrays

- Comparing two arrays is not as simple as it may seem. Let's check whether two matrices are close to each other:

```
A = np.array([0.,0.])
B = np.array([0.,0.])
if abs(B-A) < 1e-10:
    print("The two arrays are close enough")
```

ValueError: The truth value of an array with more than one element is ambiguous. Use `a.any()` or `a.all()`

- The reason why the if statement lead error message is the output of that comparison return a boolean array. It is simply an array for which the entries have the type `bool`. For example:

```
TEST = np.array([True,False])
print(TEST.dtype)
```

`bool`

Boolean array

- Any comparison operator acting on arrays will create a Boolean array instead of a simple Boolean

```
M = np.array([[2, 3], [1, 4]])  
print(M > 2)  
print(M == 0)  
N = np.array([[2, 3], [0, 0]])  
print(M == N)
```

```
[[False  True]  
 [False  True]  
 [  
  ]  
 [  
  ]
```

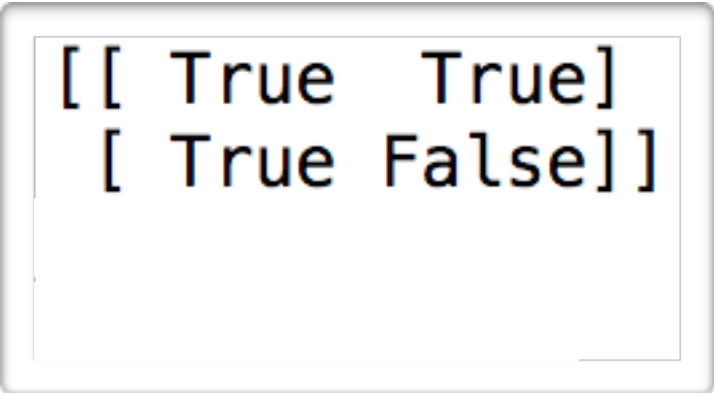
- So what is the outcome of the follows?

```
print(abs(B-A) < 1e-10)
```

Functions `all` and `any`

- Therefore, one cannot use array comparison directly in conditional statements, for example, `if` statements. the solution is to use the methods `all` and `any`

```
A = np.array([[1,2],[3,4]])
B = np.array([[1,2],[3,3]])
print(A == B)
print((A == B).all())
print((A != B).any())
if (abs(B-A) < 1e-10).all():
    print("The two arrays are close enough")
```



```
[[ True  True]
 [ True False]]
```

Checking for equality

- Note that two floats may be very close without being equal. In NumPy, it is possible to check for equality by function `allclose` with a given precision.

```
data = np.random.rand(2)*1e-3
small_error = np.random.rand(2)*1e-8
data_err = data + small_error
print(data)
print(data_err)
print(data == data_err)
print(np.allclose(data, data_err, rtol=1.e-5, atol=1.e-8))
```

```
[0.00076381 0.00017465]
[0.00076382 0.00017465]
[False False]
True
```

- The tolerance is given in terms of a relative tolerance bound, `rtol`, and an absolute error bound, `atol`. The command `allclose` is a short form of:

```
(abs(A-B) < atol+rtol*abs(B)).all()
```

- Note that `allclose` can be also applied to scalars.

Boolean operations on arrays

- You cannot use `and`, `or`, and `not` on Boolean arrays. Instead, we can use these operators for **componentwise** logical operations on Boolean arrays:

Logic operator	Replacement for Boolean arrays
A and B	A & B
A or B	A B
not A	~ A

```
A = np.array([True, True, False, False])
B = np.array([True, False, True, False])
print(A & B)
print(A | B)
print(~A)
print(A and B)
```


```
[
[
[
```

ValueError: The truth value of an array with more than one element is ambiguous. Use `a.any()` or `a.all()`

Indexing with Boolean arrays

- It is often useful to access and modify only **parts of an array**, depending on its value. For instance, one might want to access all the positive elements of an array. This turns out to be possible using Boolean arrays, which act like **masks** to select only some elements of an array.
- The result of such an indexing is always a **vector**.

```
B = np.array([[True, False], [False, True]])
M = np.array([[2, 3], [1, 4]])
print(M[B])
M[B] = 0
print(M)
M[B] = 10, 20
print(M)
M[M>2] = 0
print(M)
```



When will we use Boolean array Indexing?

- Now we have a sequence of data with some measurement error. Suppose further that we run a regression and it gives us a deviation for each value. We wish to obtain
 1. all the exceptional values with absolute error > 0.5
 2. all the good values which are lower than 5.0

```
data = np.linspace(1,10,10) # data
error = np.random.rand(10) # the errors
exceptional = data[abs(error) > 0.5]
small = data[(abs(error) <= 0.5) & (data < 5.)]
print(data)
print(error)
print(exceptional)
print(small)
```

```
[ 1.  2.  3.  4.  5.  6.  7.  8.  9. 10.]
[0.03212591 0.87259363 0.21471738 0.79291333 0.33642996 0.2386707
 0.13483616 0.73206135 0.3556771 0.90487261]
[ 2.  4.  8. 10.]
[1. 3.]
```

Using where

- The command `where(condition, a, b)` can take a Boolean array as a condition. This will return values from `a` when the condition is `True` and values from `b` when it is `False`.
- For instances consider, a Heaviside function:

```
def H(x):  
    return np.where(x < 0, 0, 1)
```

```
x = np.linspace(-1,1,11)  
print(x)  
print(H(x))
```

$$H(x) = \begin{cases} 0 & x < 0 \\ 1 & x \geq 0 \end{cases}$$

```
[-1.  -0.8 -0.6 -0.4 -0.2  0.   0.2  0.4  0.6  0.8  1. ]  
[0  0  0  0  0  1  1  1  1  1  1]
```

Using where

- More example to demonstrated how to manipulate elements from an array or a scalar depending on a condition:

```
x = np.linspace(-4,4,5)
print(x)
print(np.where(x > 0, x**2+3, 0))
print(np.where(x > 0, 1, -1))
```

```
[-4. -2.  0.  2.  4.]
[ 0.  0.  0.  7. 19.]
[-1 -1 -1  1  1]
```

- If the second and third arguments are omitted, then a **tuple** containing the **indexes** of the elements satisfying the condition is returned.

```
a = np.linspace(-4,4,9)
b = a.reshape((3,3))
print(a)
print(np.where(a > 0 ))
print(b)
print(np.where(b > 0))
```


```
[-4. -3. -2. -1.  0.  1.  2.  3.  4.]

[[-4. -3. -2.]
 [-1.  0.  1.]
 [ 2.  3.  4.]]
(array([1, 2, 2, 2]), array([ ]))
```

Broadcasting


- Broadcasting in NumPy denotes the ability to guess a common, compatible shape between two arrays.

```
vector = np.arange(4)
vector2 = vector + 1.
vector[0]=3.3
print(vector)
print(vector2)
```



- In this example, everything happens as if the scalar 1. had been converted to an array of the same length as `vector`, that is, `array([1., 1., 1., 1.])`, and then added to `vector`.

- ```
C = np.arange(1,3).reshape(-1,1) # column
R = np.arange(2).reshape(1,-1) # row
print(C)
print(R)
print(C + R)
```



# Quiz

|    |    |    |   |   |   |   |   |    |    |    |   |   |   |   |   |    |    |    |
|----|----|----|---|---|---|---|---|----|----|----|---|---|---|---|---|----|----|----|
| 0  | 0  | 0  |   | 0 | 1 | 2 |   | 0  | 0  | 0  |   | 0 | 1 | 2 |   | 0  | 1  | 2  |
| 10 | 10 | 10 | + |   |   |   | = | 10 | 10 | 10 | + | 0 | 1 | 2 | = | 10 | 11 | 12 |
| 20 | 20 | 20 |   |   |   |   |   | 20 | 20 | 20 |   | 0 | 1 | 2 |   | 20 | 21 | 22 |
| 30 | 30 | 30 |   |   |   |   |   | 30 | 30 | 30 |   | 0 | 1 | 2 |   | 30 | 31 | 32 |

## Shape mismatch

4x3

|    |    |    |
|----|----|----|
| 0  | 0  | 0  |
| 10 | 10 | 10 |
| 20 | 20 | 20 |
| 30 | 30 | 30 |

+

4

|   |   |   |   |
|---|---|---|---|
| 0 | 1 | 2 | 3 |
|---|---|---|---|

mismatch!

ValueError:

together wi

**ValueError:** operands could not be broadcast together with shapes (4,3) (4,)



Solve a Linear Algebra by Python

# Linear algebra operations

- The essential operator that performs most of the usual operations of linear algebra is the Python function `dot`. It is used for computing matrix-vector multiplications, a scalar product between two vectors, matrix-matrix products.

```
A = np.array([[1., 2, 3], [1, 1, 3], [1, 2, 4]])
b = np.array([10., 8, 13])
M = np.array([[1., 2, 3], [4, 5, 6], [7, 8, 9]])
print(A)
print(b)
print(np.dot(A, b)) # Matrix dot vector
print(np.dot(A, M)) # Matrix dot Matrix
```

```
[[1. 2. 3.]
 [1. 1. 3.]
 [1. 2. 4.]]
[10. 8. 13.]
[65. 57. 78.]
[[30. 36. 42.]
 [26. 31. 36.]
 [37. 44. 51.]]
```

# Solving a linear system

- If  $A$  is a matrix and  $b$  is a vector, you can solve the linear equation:  $Ax = b$
- Using the `solve` method in scipy linear algebra

In previous page,  $A$  and  $b$

```
[[1. 2. 3.]
 [1. 1. 3.]
 [1. 2. 4.]
 [10. 8. 13.]
```

```
import scipy.linalg as sl
```

```
x = sl.solve(A, b)
print(x)
```

```
[-3. 2. 3.]
[0. 0. 0.] 0.0
True
True
```

```
print(np.dot(A, x)-b, sl.norm(np.dot(A, x)-b))
print(sl.norm(np.dot(A, x)-b)<1e-6) # old school
print(np.allclose(np.dot(A, x), b))
```

- The command `norm` is the positive length to each vector.

$$\|x\|_2 := \sqrt{x_1^2 + \cdots + x_n^2}.$$

- The command `allclose` is used here to compare two vectors. If they are close enough to each other, this command returns `True`. The default tolerance value is  $10^{-8}$ .

# More methods in `scipy.linalg`

| Methods                      | Description (matrix methods)                               |
|------------------------------|------------------------------------------------------------|
| <code>sl.det</code>          | Determinant of a matrix                                    |
| <code>sl.eig</code>          | Eigenvalues and eigenvectors of a matrix                   |
| <code>sl.inv</code>          | Matrix inverse                                             |
| <code>sl.pinv</code>         | Matrix pseudoinverse                                       |
| <code>sl.norm</code>         | Matrix or vector norm                                      |
| <code>sl.svd</code>          | Singular value decomposition                               |
| <code>sl.lu</code>           | LU decomposition                                           |
| <code>sl.qr</code>           | QR decomposition                                           |
| <code>sl.cholesky</code>     | Cholesky decomposition                                     |
| <code>sl.solve</code>        | Solution of a general or symmetric linear system: $Ax = b$ |
| <code>sl.solve.banded</code> | The same for banded matrices                               |
| <code>sl.lstsq</code>        | Least squares solution                                     |

# Methods: Examples

```
A = np.array([[1., 2, 3], [1, 1, 3], [1, 2, 4]])
b = np.array([10., 8, 13])
```

```
InvA = sl.inv(A)
print(InvA)
print(np.dot(InvA, b))
print(sl.solve(A, b))

print(sl.det(A))
(eigVal, eigVec) = sl.eig(A)
print(eigVal)
print(eigVec)
```

$$\begin{cases} x+2y+3z=10 \cdots (1) \\ x+ y+3z=8 \cdots (2) \\ x+2y+4z=13 \cdots (3) \end{cases}$$

```
[[2. 2. -3.]
 [1. -1. 0.]
 [-1. 0. 1.]]
[-3. 2. 3.]
[-3. 2. 3.]
```

```
-0.999999999999999998
[6.29257994+0.j 0.2783488 +0.j -0.57092874+0.j]
[[-0.56270964 -0.91432584 0.58695504]
 [-0.4855477 -0.19908612 -0.78101129]
 [-0.66903012 0.35266553 0.21331935]]
```

# Methods: Examples

- `inv()` and `solve()` can only be applied to a square matrix. `pinv()`, however, can determine the inverse matrix numerically.

```
A = np.array([[1., 2, -1], [0, -5, 3]])
b = np.array([4., 1])
InvA = sl.pinv(A)
print(InvA)
x_pinv = np.dot(InvA, b)
print(x_pinv)
print(sl.norm(np.dot(A, x_pinv) - b) < 1e-6)
```

$$\begin{bmatrix} 1 & 2 & -1 \\ 0 & -5 & 3 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 4 \\ 1 \end{bmatrix}$$

```
[[0.97142857 0.37142857]
 [0.08571429 -0.11428571]
 [0.14285714 0.14285714]]
[4.25714286 0.22857143 0.71428571]
True
```

```
InvA = sl.inv(A)
x = sl.solve(A, b)
```


**ValueError:** expected square matrix

**ValueError:** Input a needs to be a square matrix.



# Methods: Examples

- The method of least squares is a standard approach in regression analysis to approximate the solution.
- “Least squares” means that the overall solution minimizes the sum of the squares of the residuals made in the results of every single equation. i.e. that minimizes the Euclidean 2-norm  $\|b - Ax\|^2$



```
(x_lstsq, r, rank, s) = sl.lstsq(A, b)
print(x_lstsq)
print(sl.norm(np.dot(A, x_lstsq) - b) < 1e-6)
print(r, rank, s)
```

$$\begin{bmatrix} 1 & 2 & -1 \\ 0 & -5 & 3 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 4 \\ 1 \end{bmatrix}$$

The solution  $x$  → `[4.25714286 0.22857143 0.71428571]`

The residues of  $b$  → `[]`

The rank of  $A$  → `2`

The singular value of  $A$  → `[6.25339693 0.94605857]`